Recent Advances on Traveling Salesman Problem *

Yihui He
Xi'an Jiaotong University
Xi'an, China
heyihui@stu.xjtu.edu.cn

Ming Xiang †
Xi'an Jiaotong University
Xi'an China
mxiang@mail.xjtu.edu.cn

Abstract

With applications to many disciplines, the traveling salesman problem (TSP) is a classical computer science optimization problem with applications to industrial engineering, theoretical computer science, bioinformatics, and several other disciplines. In recent years, there have been a plethora of novel approaches for approximate solutions ranging from simplistic greedy to cooperative distributed algorithms derived from artificial intelligence. In this paper, we perform an evaluation and analysis of cornerstone algorithms for the metric TSP. We evaluate the nearest neighbor, greedy, Christofides, and genetic algorithms. We use several datasets as input for the algorithms including several small datasets, two medium-sized datasets representing cities in the United States, and a synthetic dataset consisting of 1,000 cities to test algorithm scalability. We discover that the nearest neighbor and greedy algorithms efficiently calculate solutions for smaller datasets. Christofides has the best performance for both optimality and runtime for medium to large datasets. Genetic algorithms can occasionally find near-optimal solutions but have no guarantee and generally have longer runtimes.

1. Introduction

Known to be NP-hard, the traveling salesman problem (TSP) was first formulated in 1930 and is one of the most studied optimization problems to date [14]. The problem is as follows: given a list of cities and a distance between each pair of cities, find the shortest possible path that visits every city exactly once and returns to the starting city. The TSP has broad applications including: shortest-path for lasers to sculpt microprocessors and delivery logistics for mail services, to name a few.

*This work was completed as an artificial intelligence course project at Xi'an Jiaotong University. Permission to copy without all or part of this material is granted provided that the copies are not made or distributed for direct commercial advantage.

†corresponding author.

The TSP is an area of active research. In fact, several variants have been derived from the original TSP. In this paper, we focus on the metric TSP. In the metric TSP, all distances between cities satisfy the triangle inequality. That is, for three cities, A, B, and C:

$$dist(A,C) < dist(A,B) + dist(B,C) \tag{1}$$

This simplification allows us to survey several cornerstone algorithms without introducing complex scenarios, specifically for Christofides. The remainder of this paper is organized as follows. In Section 2, we briefly review the first solutions and survey modern approaches and variants to the TSP. We describe the algorithms used in our experiment and outline key implementation details in Section 3. A description of the benchmark datasets and results of the experiment are detailed in Section 4. A discussion in Section 5 explains the findings and compares the performance of the algorithms. We then conclude and describe future work in Section 6.

2. Background

An example TSP is illustrated in Figure 1. The input is shown in subfigure (a) as a collection of cities in the twodimensional space. This input can be represented as a distance matrix for each pair of cities or as a list of points denoting the coordinate of each city. In the latter method, distances are calculated using Euclidean geometry. A nonoptimal tour is shown in subfigure (b). Although not shown in the figure, each edge will have some non-negative edge weight denoting the distance between two nodes or cities. Due to the computational complexity of the TSP, it may be necessary to approximate the optimal solution. The optimal tour is shown in subfigure (c). For small graphs, it may be possible to perform an exhaustive search to obtain the optimal solution. However, as the number of cities increases, so does the solutions space, problem complexity, and running time.

Figure 2 lists the number of edges and total possible num- ber of tours for a specific dataset size. The number

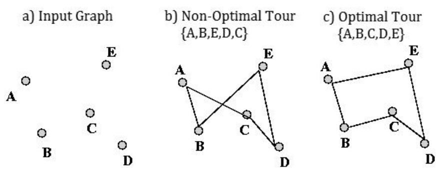

Figure 1.

of possi- ble tours is (n 1)!/2 since the same tour, with start point X and Y appears twice: once with X as the start node and once with Y as the start node.

Mathematical problems similar to the Traveling salesman problem date back to the 18th century. The basis of the problem was first discussed by Irish mathematician William Rowan Hamilton and by a British mathematician Thomas Penyngton Kirkman.

The TSP problem itself was first formulated in the 1930s by Karl Menger in Vienna and Harvard. It was later studied by statisticians Mahalanobis, Jessen, Gosh, and Marks for agricultural applications. The problem was then popularized by mathematician Merill Flood with his colleagues at Research and Development Corporation in the 1940s. By the mid-1950s, solutions for TSP began to appear. The first solution was published by Dantzig, Fulkerson, and Johnson using a dataset of 49 cities. The progression of solutions is shown in Figure 3. In 1972, Richard M. Karp proved that the Hamiltonian cycle problem was NP-Complete, which proves that the TSP is NP-Hard.

In modern day, the traveling salesman problem has a variety of applications to numerous fields. Examples among these applications include genome sequencing, air traffic con- trol, supplying manufacturing lines, and optimization.

2.1. TSP Variants

Several variants of the original TSP exist. Some of these variants differ in their representation of cities and some differ by the constraints placed on city distances. We describe the Euclidean, Graphic, and Asymmetric TSP variants in this section.

2.1.1 Euclidean TSP

In the Euclidean TSP, the input is given as a list of coordinates describing the location of each city in R 2 [1]. It is possible to extend this into higher dimensions as well. The alternative to giving a list of city coordinates is to give a distance/adjacency matrix describing the distances between each city pair. It is possible to derive the distance matrix given the city coordinates using Euclidean geometry. However, given the distance matrix, inferring the coordinates is not as straightforward and requires computation of a Gramian matrix and application of several matrix decomposition methods

2.1.2 Graphic TSP

In recent years, the graphic TSP has attracted attention from the research community. The graphic TSP problem asks for the shortest tour that visits each vertex at least once. The current best approximation algorithm returns a solution within 13/9 1.444 of the optimal [16]. For the past 30 years, Christofides had been the leading algorithm with a 1.5 approximation. Although the 13/9 approximation is for a TSP variant, it is an important milestone for TSP approximations nonetheless.

2.1.3 Asymmetric TSP

All TSP variants described thus far have assumed an undirected graph, resulting in cost(A, B) = cost(B, A). The asymmetric TSP introduces directed edges. As a consequence, algorithms with assumptions about reflexive distances may fail when cost(A, B) 6 = cost(B, A). Real world applications of this includes roads and highways specif- ically the case of one-way streets, detours, and alternate routes. Additionally when modeling vehicle energy usage, the geographical topology makes traveling uphill cheaper than downhill despite traveling between the same two points.

2.2. Related Heuristics

In this section, we briefly summarize current approaches for the metric TSP. We describe the 2-approximation algorithm and survey an advanced technique based on artificial intelligence.

2.2.1 2-Approximation Algorithm

A simple two-approximation solution can be achieved by constructing a minimum spanning tree (MST) for the input TSP graph and creating a list of vertices (no duplicates) from a pre-order walk of the MST [20]. This list of vertices becomes a Hamiltonian cycle and is a solution to the TSP. This algorithm completes in polynomial time and is is guaranteed to return a two-approximation solution (see [20] for proof).

2.2.2 Ant-Colony Approach

Classified under the umbrella of nature-inspired algorithms, the ant colony system (ACS) is a distributed swarm intelligence algorithm that has a set of cooperating agents called ants that attempt to solve the TSP. This method attempts to mimic how ants find the shortest path from their home to food in real life. In ACS, ants communicate by depositing pheromones on the graph edges as they build TSP solutions [7]. Over time, the shorter paths build larger amounts of pheromones. The solution to the TSP is the path with highest pheromone levels which visit every city.

This algorithm is able to run on both symmetric and asymmetric TSPs. For the symmetric TSP with 170 cities or less, the ACS algorithm finds the optimum solution. For the asymmetric TSP with 170 cities or less, it found a solution within 0.40algorithm approaches [7].

3. Algorithms

e now move to a discussion of the algorithms used in our evaluation. First, we describe the traditional nearest neighbor and greedy approaches in Sections 3.1 and 3.2. We then outline Christofides algorithm in Section 3.3 and then discuss the genetic algorithm in Section 3.4.

4. Greedy Algorithm

The greedy heuristic is based on Kruskals algorithm to give an approximate solution to the TSP [13]. The algorithm forms a tour of the shortest route and can be constructed if and only if: The edges of the tour must not form a cycle unless the selected number of edges is equal to the number of vertices in the graph. The selected edge (before being appended to the tour) does not increase the degree of any node to be more than 2. The algorithm begins by sorting all edges from least weight to most heavily weighted. After the edges are sorted, the least heavily-weighted edge is selected and it is added to the tour if it does not violate the above conditions. The algo- rithm continues by selecting the next least-cost edge and adding it to the tour. This process is repeated until all ver- tices can be reached by the tour. The result is a minimum spanning tree and is a solution for the TSP. The runtime for the greedy algorithm is O(n 2 log(n)) and generally returns a solution within 15-20% of the Held-Karp lower bound [17].

4.1. Genetic Algorithm

Genetic algorithms (GA) are search heuristics that attempt to mimic natural selection for many problems in optimization and artificial intelligence [4]. In a genetic algorithm, a population of candidate solutions is evolved over time towards better solutions. These evolutions generally occur through mutations, randomization, and recombination. We define a fitness function to differentiate between better and worse solutions. Solutions, or individuals, with higher fitness scores are more likely to survive over time. The final solution is found if the population converges to a solution within some threshold. However, great care must be taken to avoid being trapped at local optima.

We will now apply a genetic algorithm to the traveling salesman problem [2]. We define a fitness function F as the length of the tour. Supposed we have an ordering of the cities A = x 1 , x 2 , ..., x n where n is the number of cities. The fitness score for the TSP becomes the cost of the tour d(x, y) denote the distance from x to y.

$$F(A) = \sum_{i=0}^{n-1} d(x_i, x_{i+1}) + d(x_n, x_0) \tag{2}$$

The genetic algorithm begins with an initial, P 0 , random population of candidate solutions. That is, we have a set of paths that may or may not be good solutions. We then move forward one time step. During this time step, we perform a set of probabilistic and statistical methods to select, mutate, and produce an offspring population, P 1 , with traits similar to those of the best individuals (with the highest fitness) from P 0 . We then repeat this process until our population becomes homogeneous.

The running time of genetic algorithms is variable and de- pendent on the problem and heuristics used. However, for each individual in the population, we require O(n) space for storage of the path. For genetic crossover, the space require- ment remains O(n). The best genetic algorithms can find solutions within 2[10].

5. Implementation

5.1. Greedy Algorithm

The greedy solution to TSP differs from the nearest neigh- bor heuristic because it uses a Kruskals approach to the problem. Instead of starting at a random node and building a tour using the nearest neighbor of the selected node, the Greedy algorithm selects the least weighted edge and adds it to the tour regardless of if it is connected or disconnected to the current tour.

The Greedy algorithm was implemented in Java and is available on Github 1 . The program reads the input datasets from the files and constructs a distance matrix corresponding to distances between cities. The algorithm first sorts all of the weights of the edges from lowest to most heavily weighted, and selects the lightest edge to begin the tour with. It then selects the next lightest edge and adds it to the tour as long as it doesnt create a cycle or make any vertices have a degree of more than 2. The algorithm keeps performing the loop until the number of edges in the tour is equal to the number of vertices in the graph. It then prints out all of the edges added to the tour, the running time of the operation, and returns the path cost of the tour.

References

[1] A. Alpher. Frobnication. *Journal of Foo*, 12(1):234–778, 2002.

[2] A. Alpher and J. P. N. Fotheringham-Smythe. Frobnication revisited. *Journal of Foo*, 13(1):234–778, 2003.

[3] A. Alpher, J. P. N. Fotheringham-Smythe, and G. Gamow. Can a machine frobnicate? *Journal of Foo*, 14(1):234–778, 2004.

[4] Authors. The frobnicatable foo filter, 2014. Face and Gesture submission ID 324. Supplied as additional material `fg324.pdf`.

[5] Authors. Frobnication tutorial, 2014. Supplied as additional material `tr.pdf`.

[6] F. Perazzi, J. P.-T. B. McWilliams, L. Van Gool, M. Gross, and A. Sorkine-Hornung. A benchmark dataset and evaluation methodology for video object segmentation.